


# Performance of synchronous and asynchronous I/O in network applications

parallel 2015  
Hubert Schmid



# Hype Cycle

What performance can you expect from synchronous and asynchronous I/O in network applications?

# Agenda

Introduction


Example, Benchmark,  
Results

Strengths, Weaknesses,  
Hybrids

Summary



# Parallelism



optimisation of throughput time

# Concurrency



optimisation of efficiency

# Terminology

## Synchronous Model:

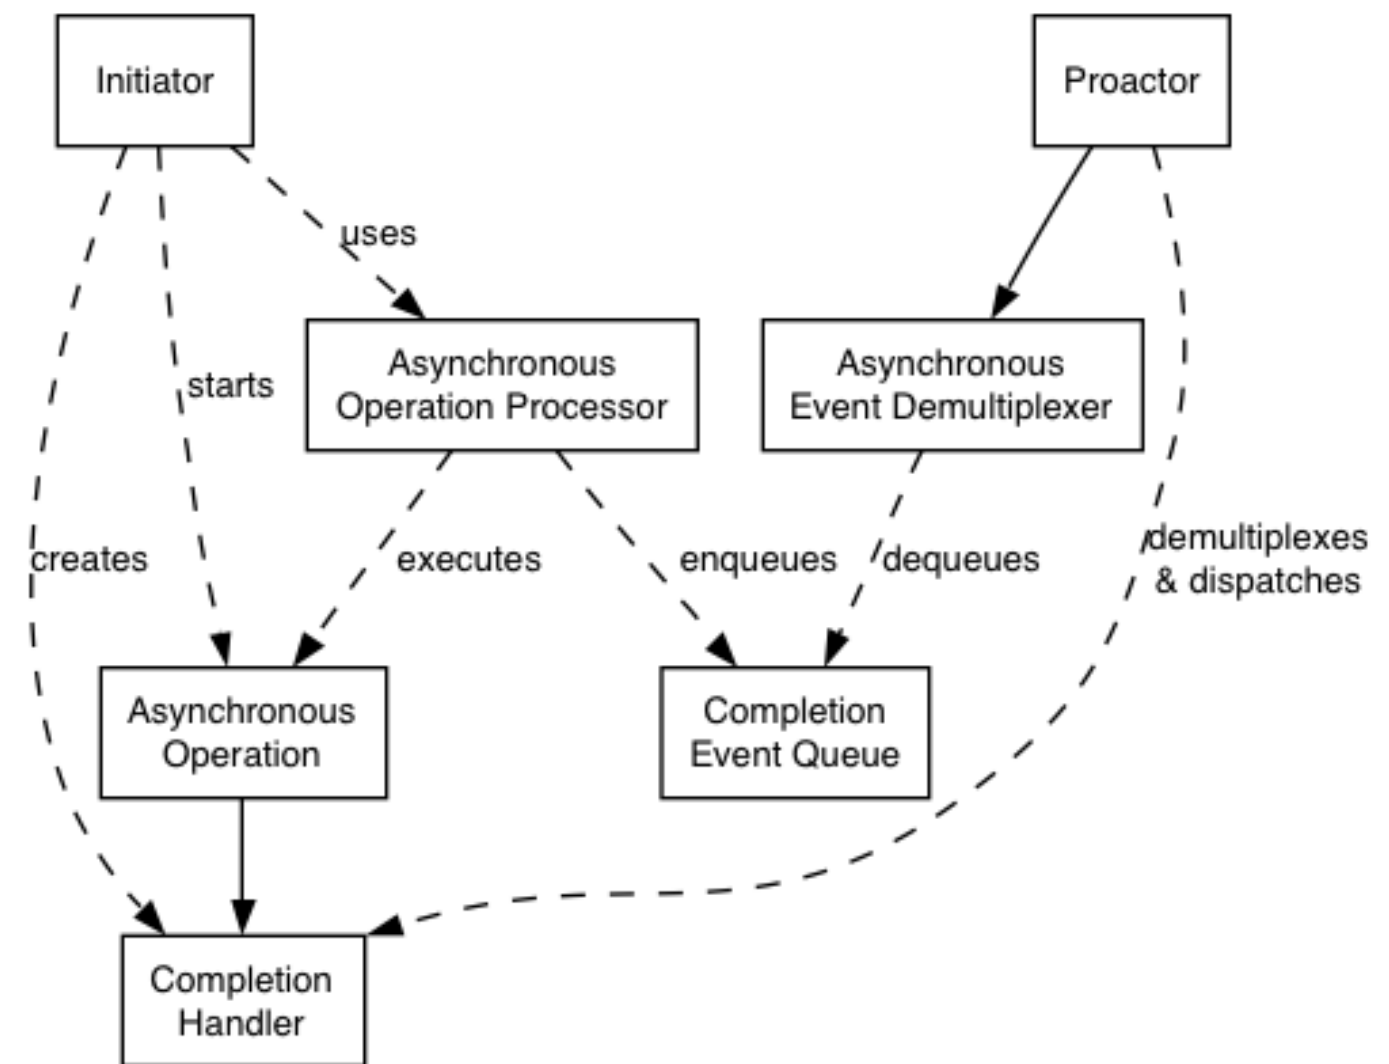
- blocking I/O
- multi-threaded

## Asynchronous Model:

- non-blocking I/O
- proactor pattern
  - notify-on-completion
  - notify-on-readiness

## Out of scope:

- asynchronism in kernel, hardware, network



Proactor Design Pattern  
(Documentation of Boost.Asio 1.57)

# Example, Benchmark, Results

# Reverse-Echo-Server



- tangible
- simple
- effective
- low-level
- realistic

# Synchronous Implementation

```
try {  
    auto timeout = 300s;  
    deadline deadline(timeout);  
    while (auto n = _stream.getline(deadline)) {  
        auto data = _stream.data();  
        std::reverse(data, data + n - 1);  
        _stream.write_n(data, n, deadline);  
        _stream.drain(n);  
        deadline.expires_from_now(timeout);  
    }  
    if (_stream.available())  
        throw std::runtime_error("data available after deadline");  
} catch (...) {  
    _handle_error();  
}
```

```
::recv(_sock, ...);  
if (errno == EAGAIN) {  
    ::poll(_sock, _timer);  
    ::recv(_sock, ...);  
}
```

```
::send(_sock, ...);  
if (errno == EAGAIN) {  
    ::poll(_sock, _timer);  
    ::send(_sock, ...);  
}
```




# Asynchronous Implementation

```
_stream.expires_from_now(300s, self);
_stream.async_getline(
    [this, self=std::move(self)](..., size_t n) {
        if (_stream.good(ec)) {
            auto data = _stream.data();
            std::reverse(data, data + n - 1);
            _stream.async_write_n(data, n,
                [this, self=std::move(self)](...) {
                    if (_stream.good(ec)) {
                        _stream.drain(n);
                        _async_run(std::move(self));
                    } else {
                        _handle_error(ec, "sending");
                    }
                });
        } else {
            _handle_error(ec, "receiving");
        }
    });
```

# Test candidates and setup

- Async Single-Core (single-threaded)
- Async Multi-Core (thread per core)
- Sync Multi-Core (thread per connection)

Measurement from 2 systems with each 500,000 simultaneous TCP connections



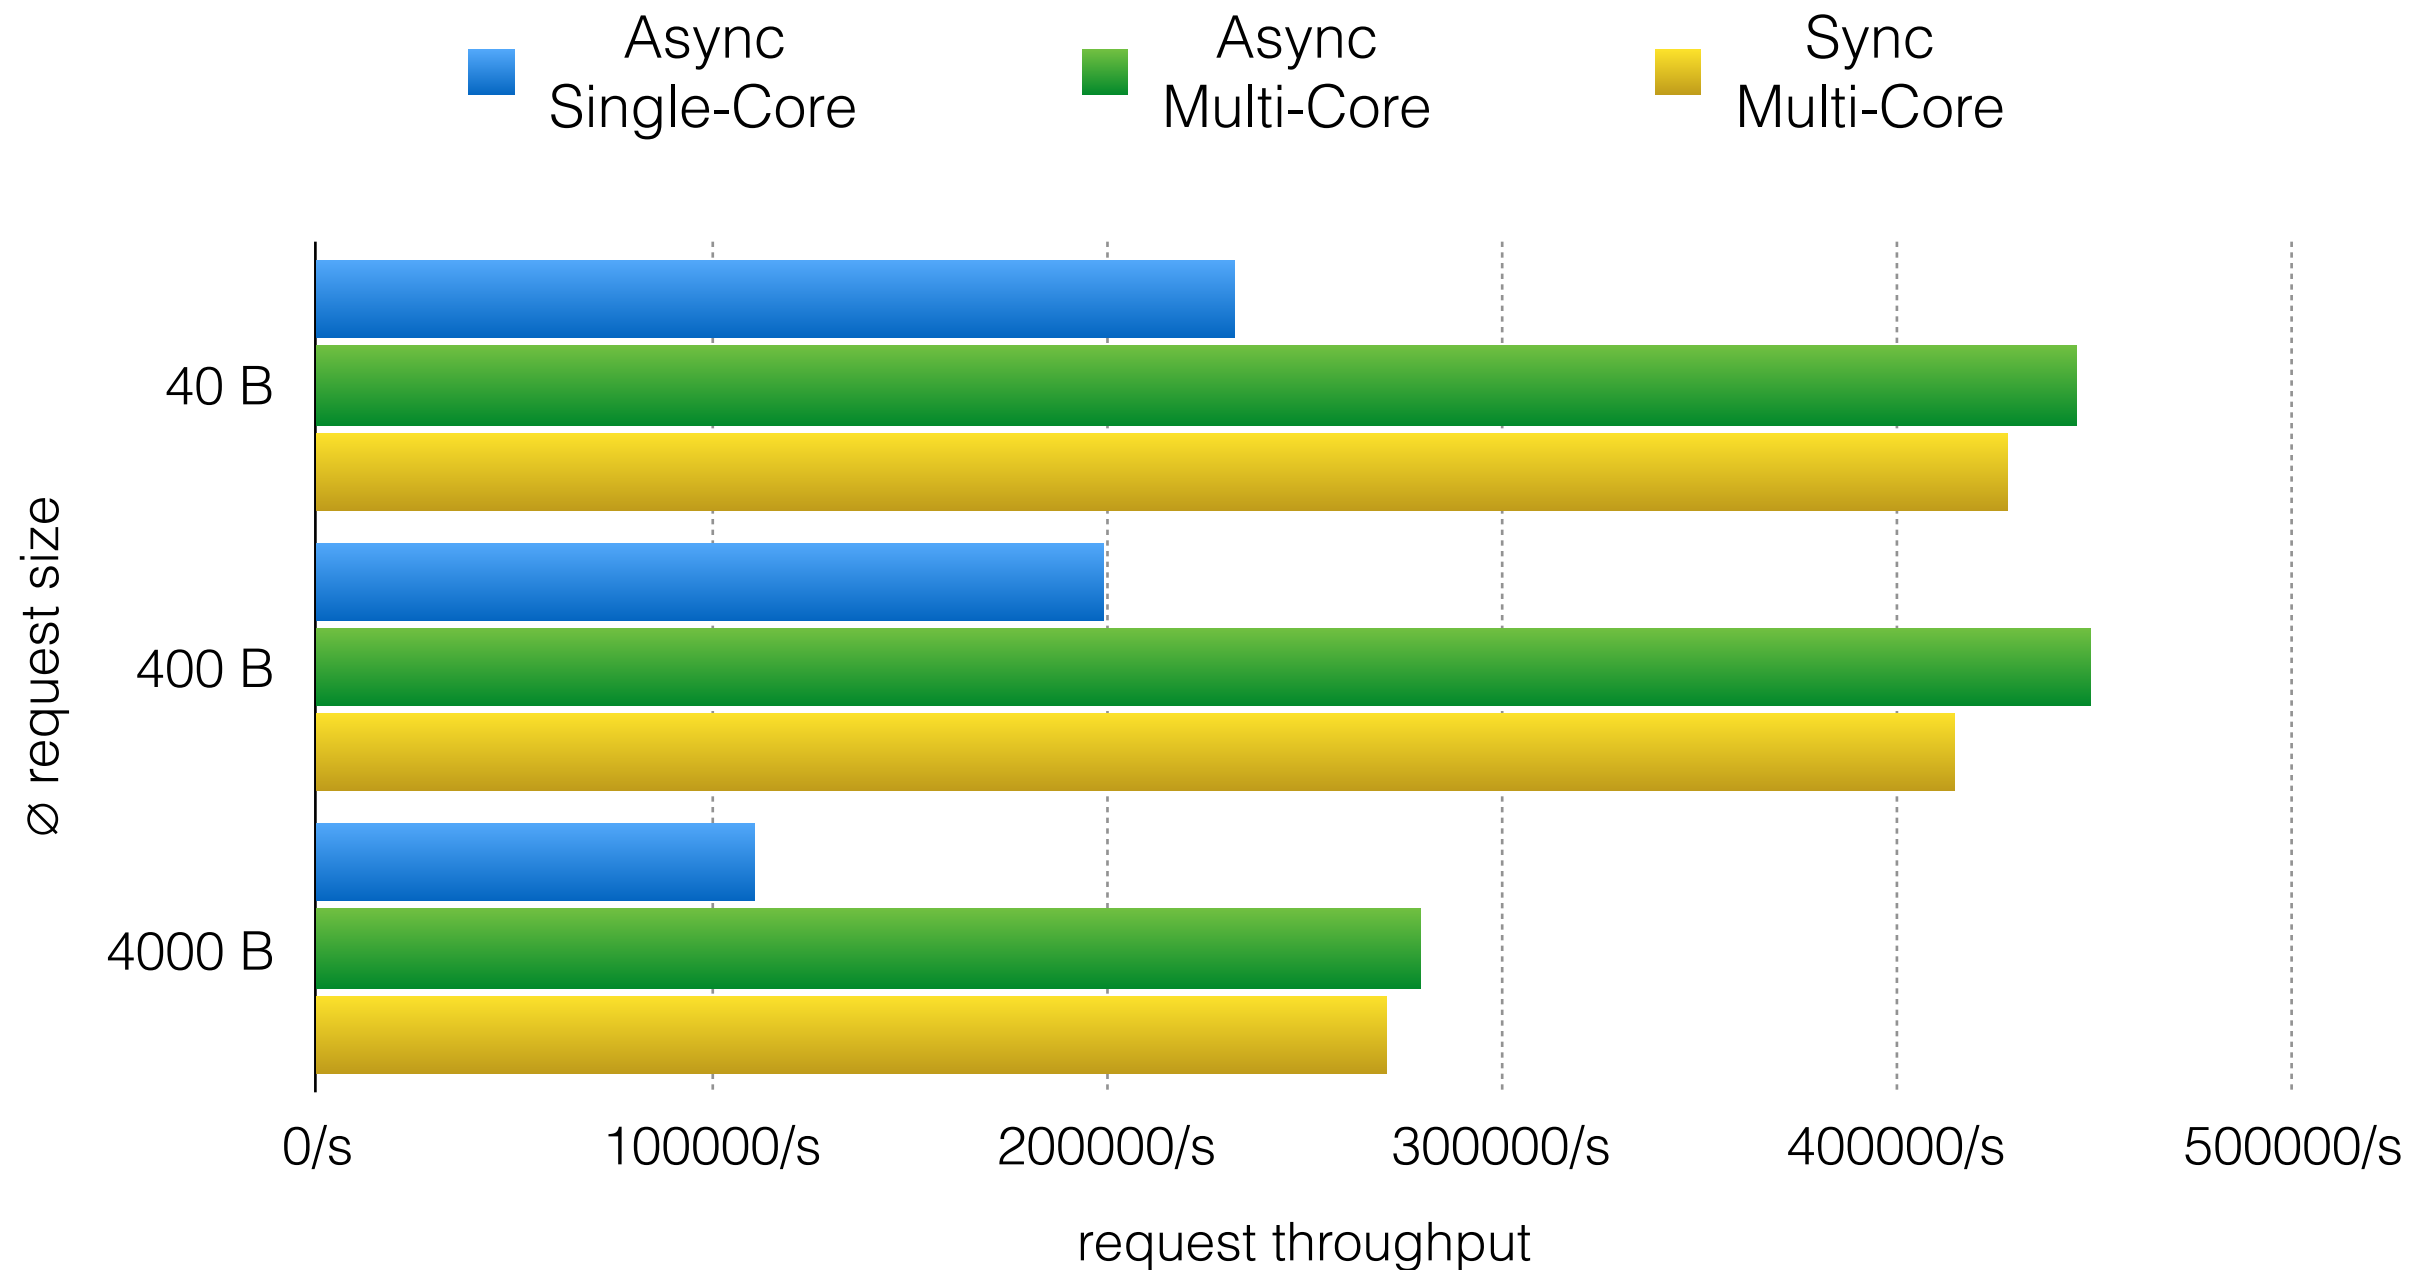
## Systems

- 2x Intel Xeon E5-2666 10C
- Intel NIC 10 GbE (2x Multi-Queue)
- Linux 3.16 (Debian jessie/testing)


## Network

- 10 Gb Ethernet (shared)

# Throughput

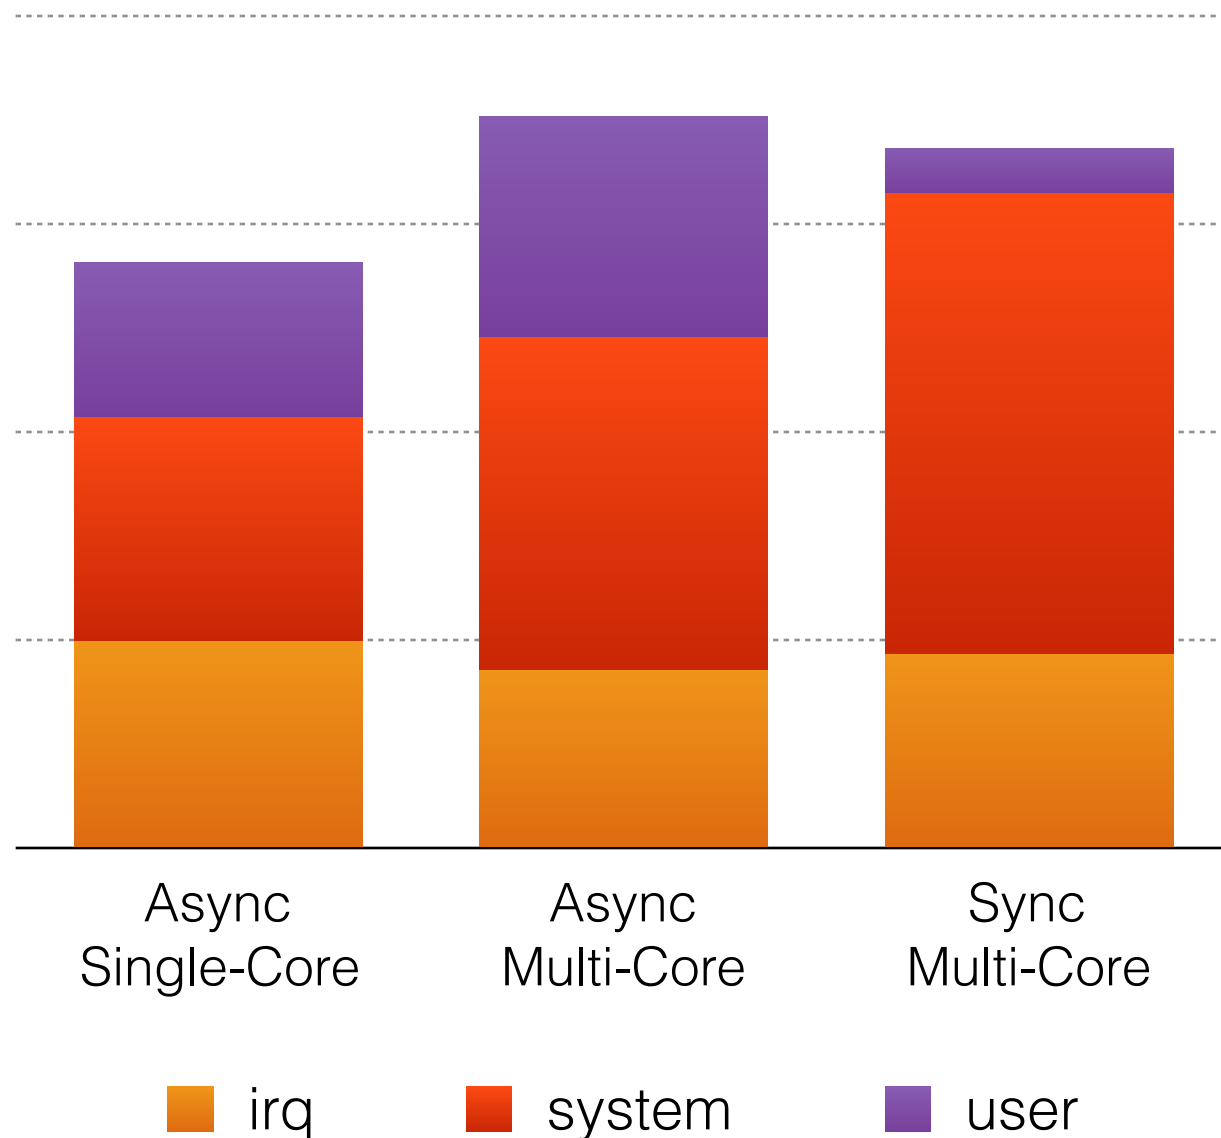


# Latency




# Resource Usage

CPU



RAM



# synchronous

```
try {
    auto timeout = 300s;
    deadline deadline(timeout);
    while (auto n = _stream.getline(
        auto data = _stream.data();
        std::reverse(data, data + n);
        _stream.write_n(data, n, dead
        _stream.drain(n);
        deadline.expires_from_now(tim
    ) {
        if (_stream.available()) {
            throw std::runtime_error("pro
        }
    }
    catch (...) {
        _handle_error();
    }
```

# system calls

# thread context switch

# asynchronous

```
_stream.expires_from_now(300s, se
_stream.async_getline(
    [this, self=std::move(self)](...)
    if (ec) {
        auto data = _stream.data();
        std::reverse(data, data + n
        _stream.async_write_n(data,
        [this, self=std::move(self
        if (ec) {
            _stream.drain(n);
            _async_run(std::move(
        } else {
            _handle_error(ec, "se
        }
    });
} else {
    _handle_error(ec, "receivin
```

# task context switch

# Results

- network saturation with multi-core
- similar usage of network stack
- different RAM usage
- actual optimisations are unrelated to programming model (e.g. receive-side-scaling with NIC support)

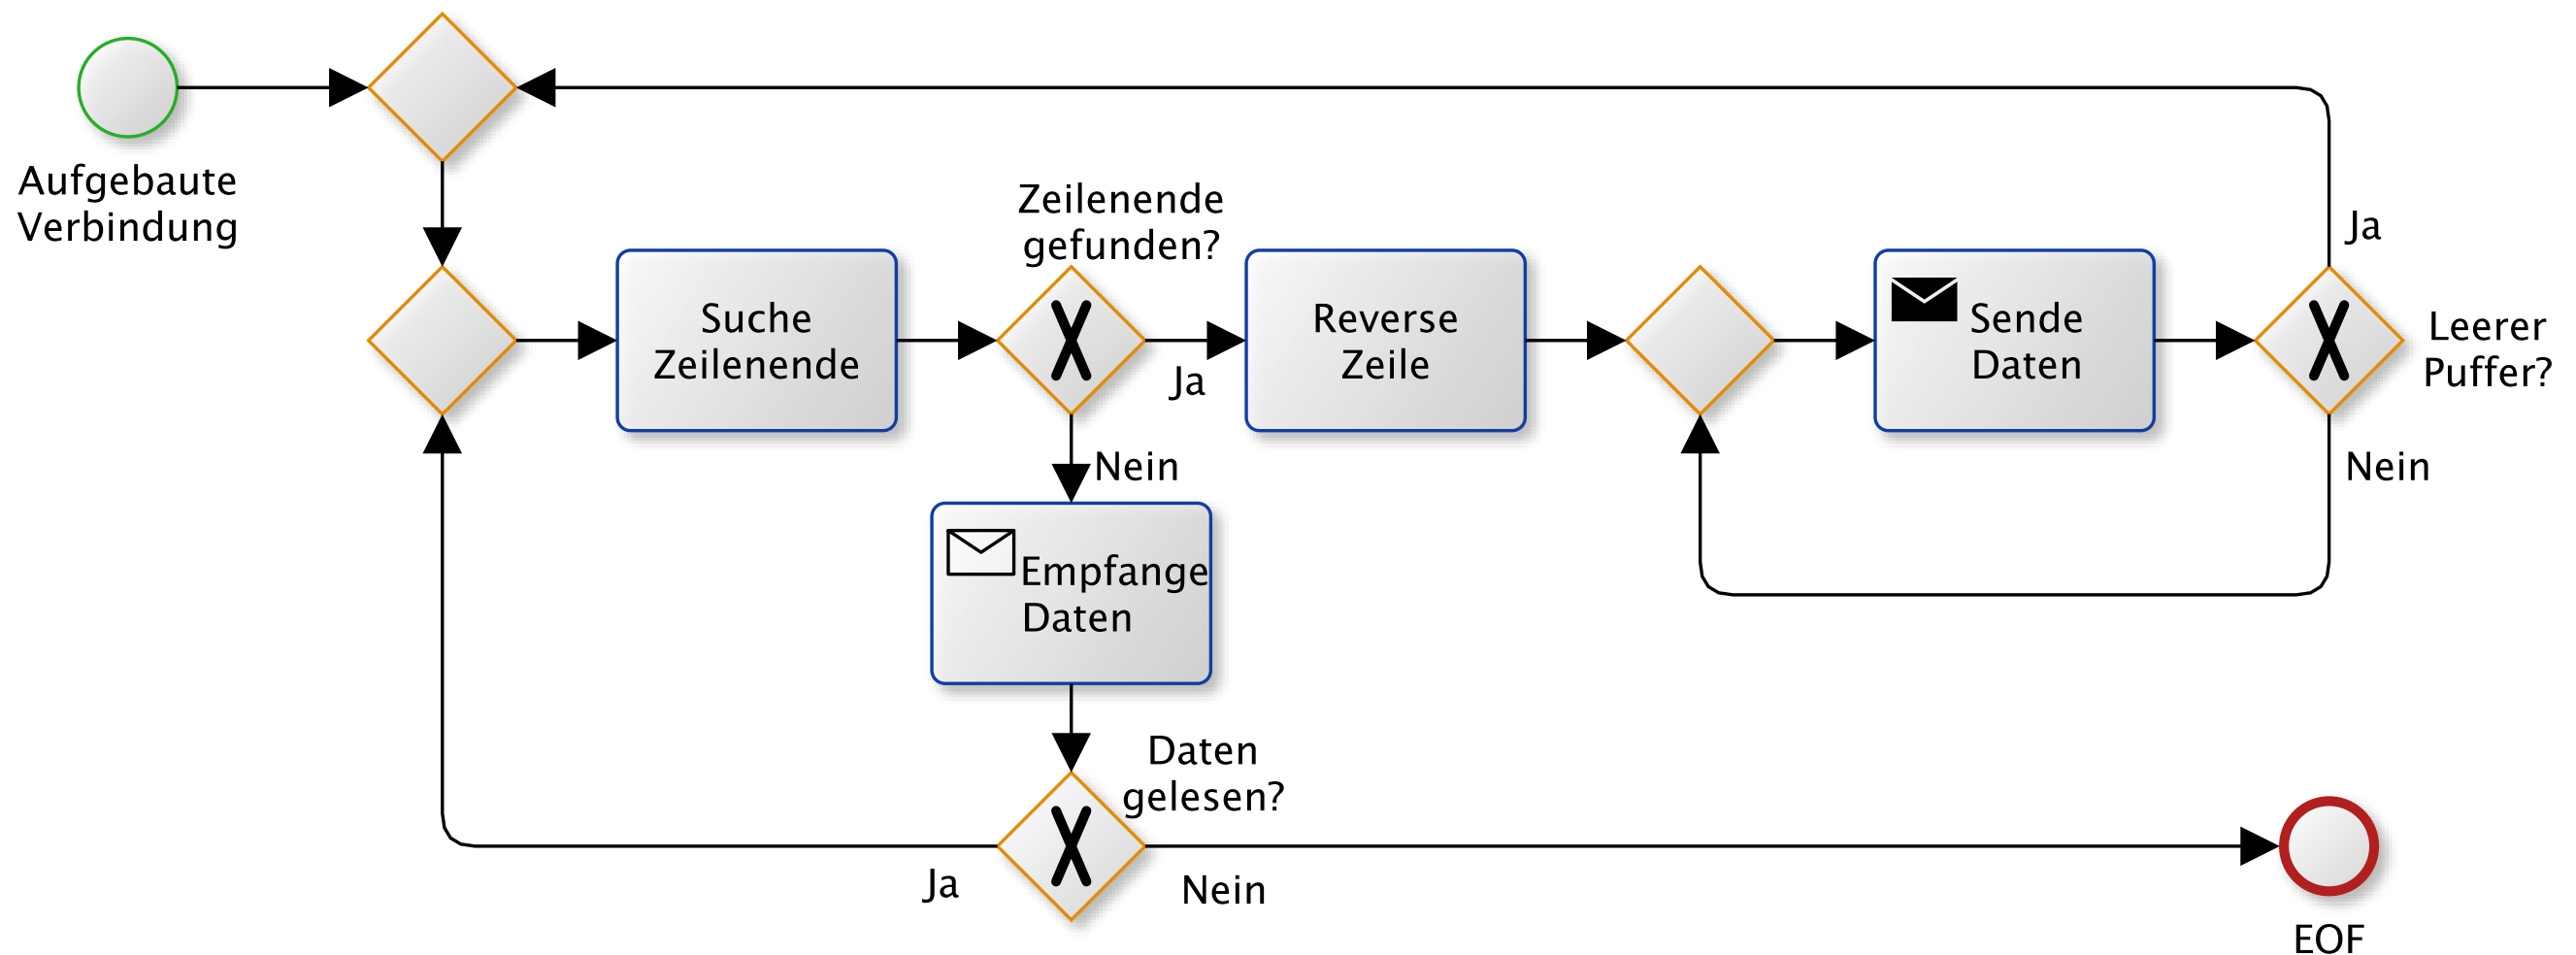
Strengths,  
Weaknesses, Hybrids



# Asynchronous Strengths



- platform support
- kernel bypass
- customisation
- non-sequential flows

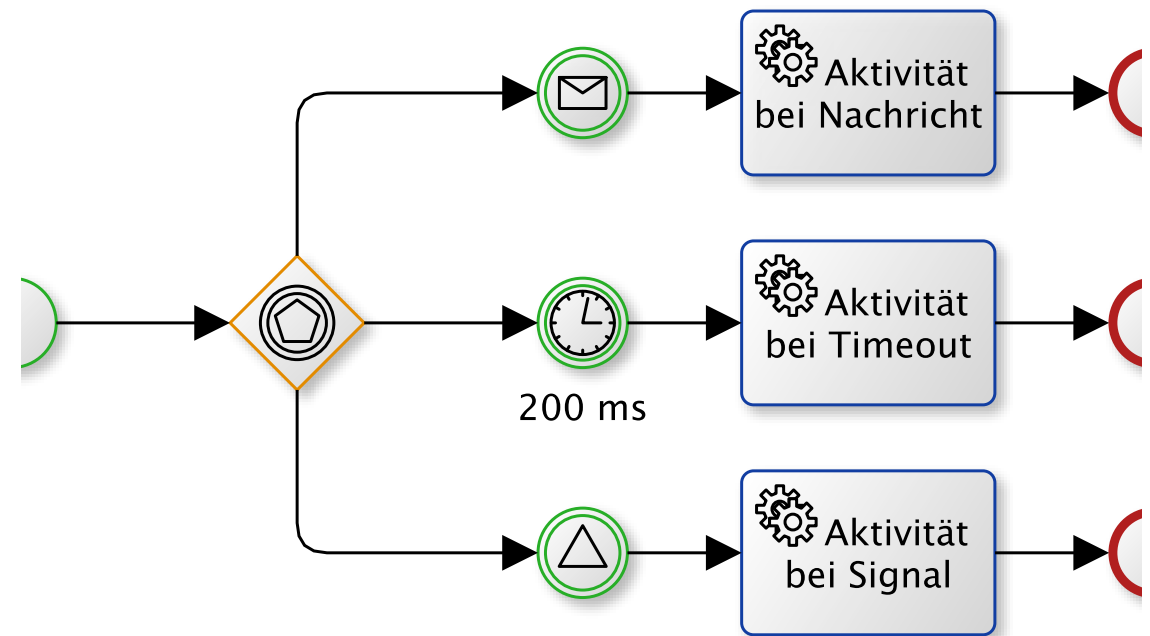



# sequence flow diagram

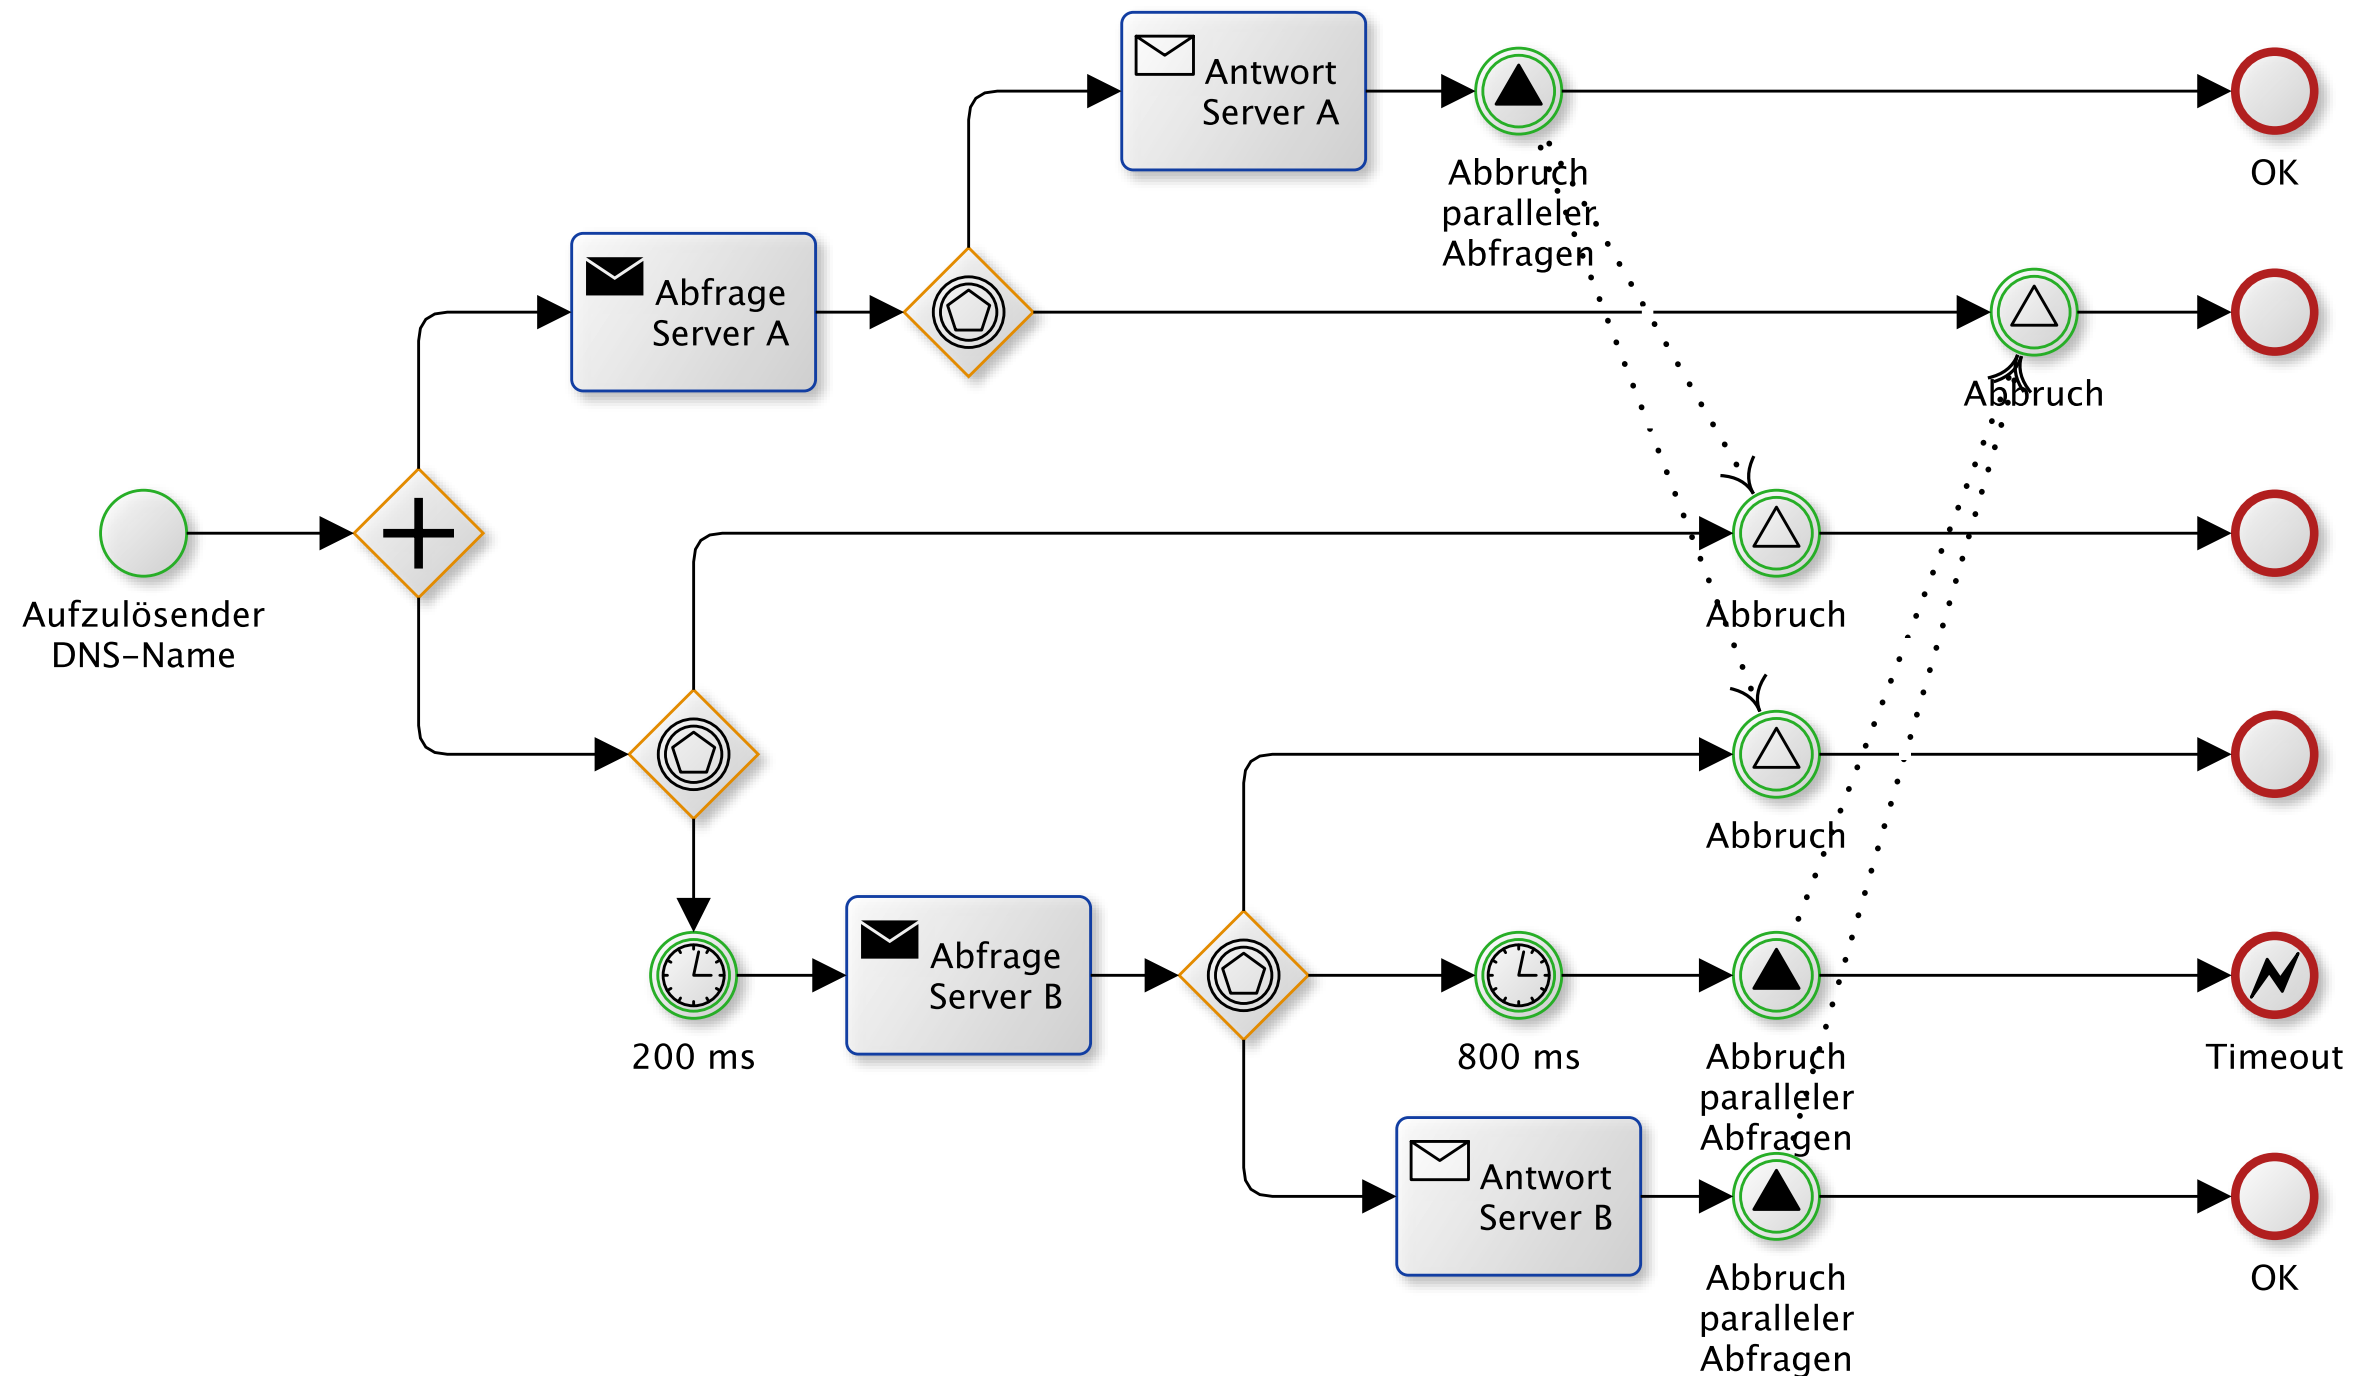
Reverse-Echo-Server is fully sequential

# Non-sequential flows


- parallel forks (overlapping I/O)
- event-based decision gateways
- control of parallel execution paths



# Example DNS Query



# Asynchronous Weaknesses



## **operating systems**

- incomplete interface (kernel, core libraries)
- virtual memory

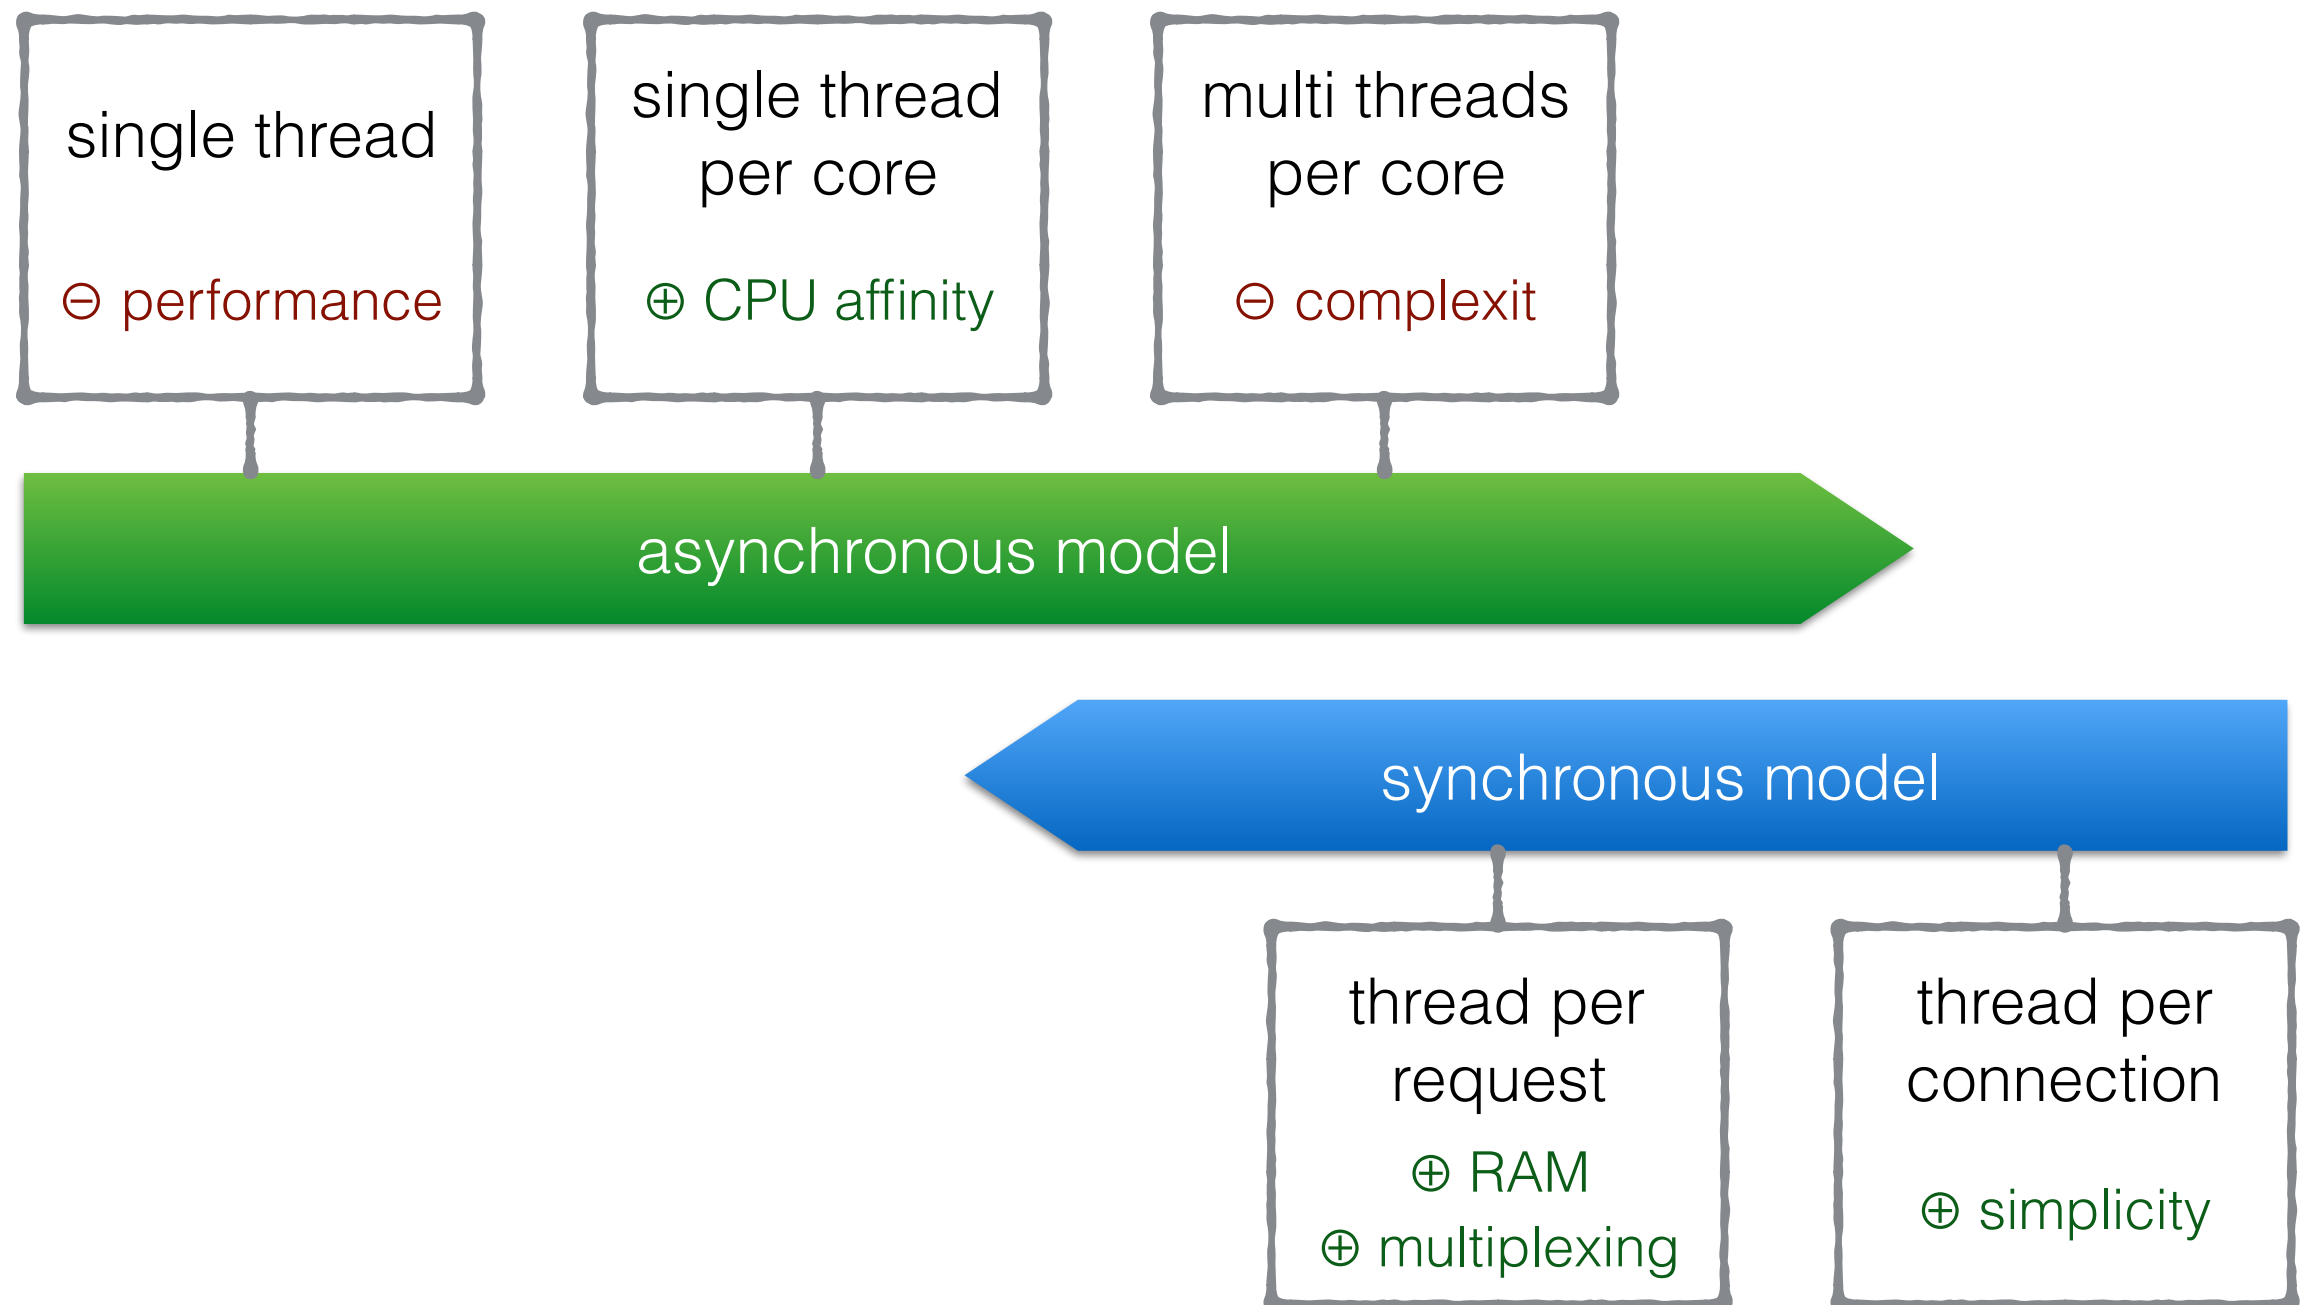
## **CPU-bound flows**

- no preemption (causing latency)

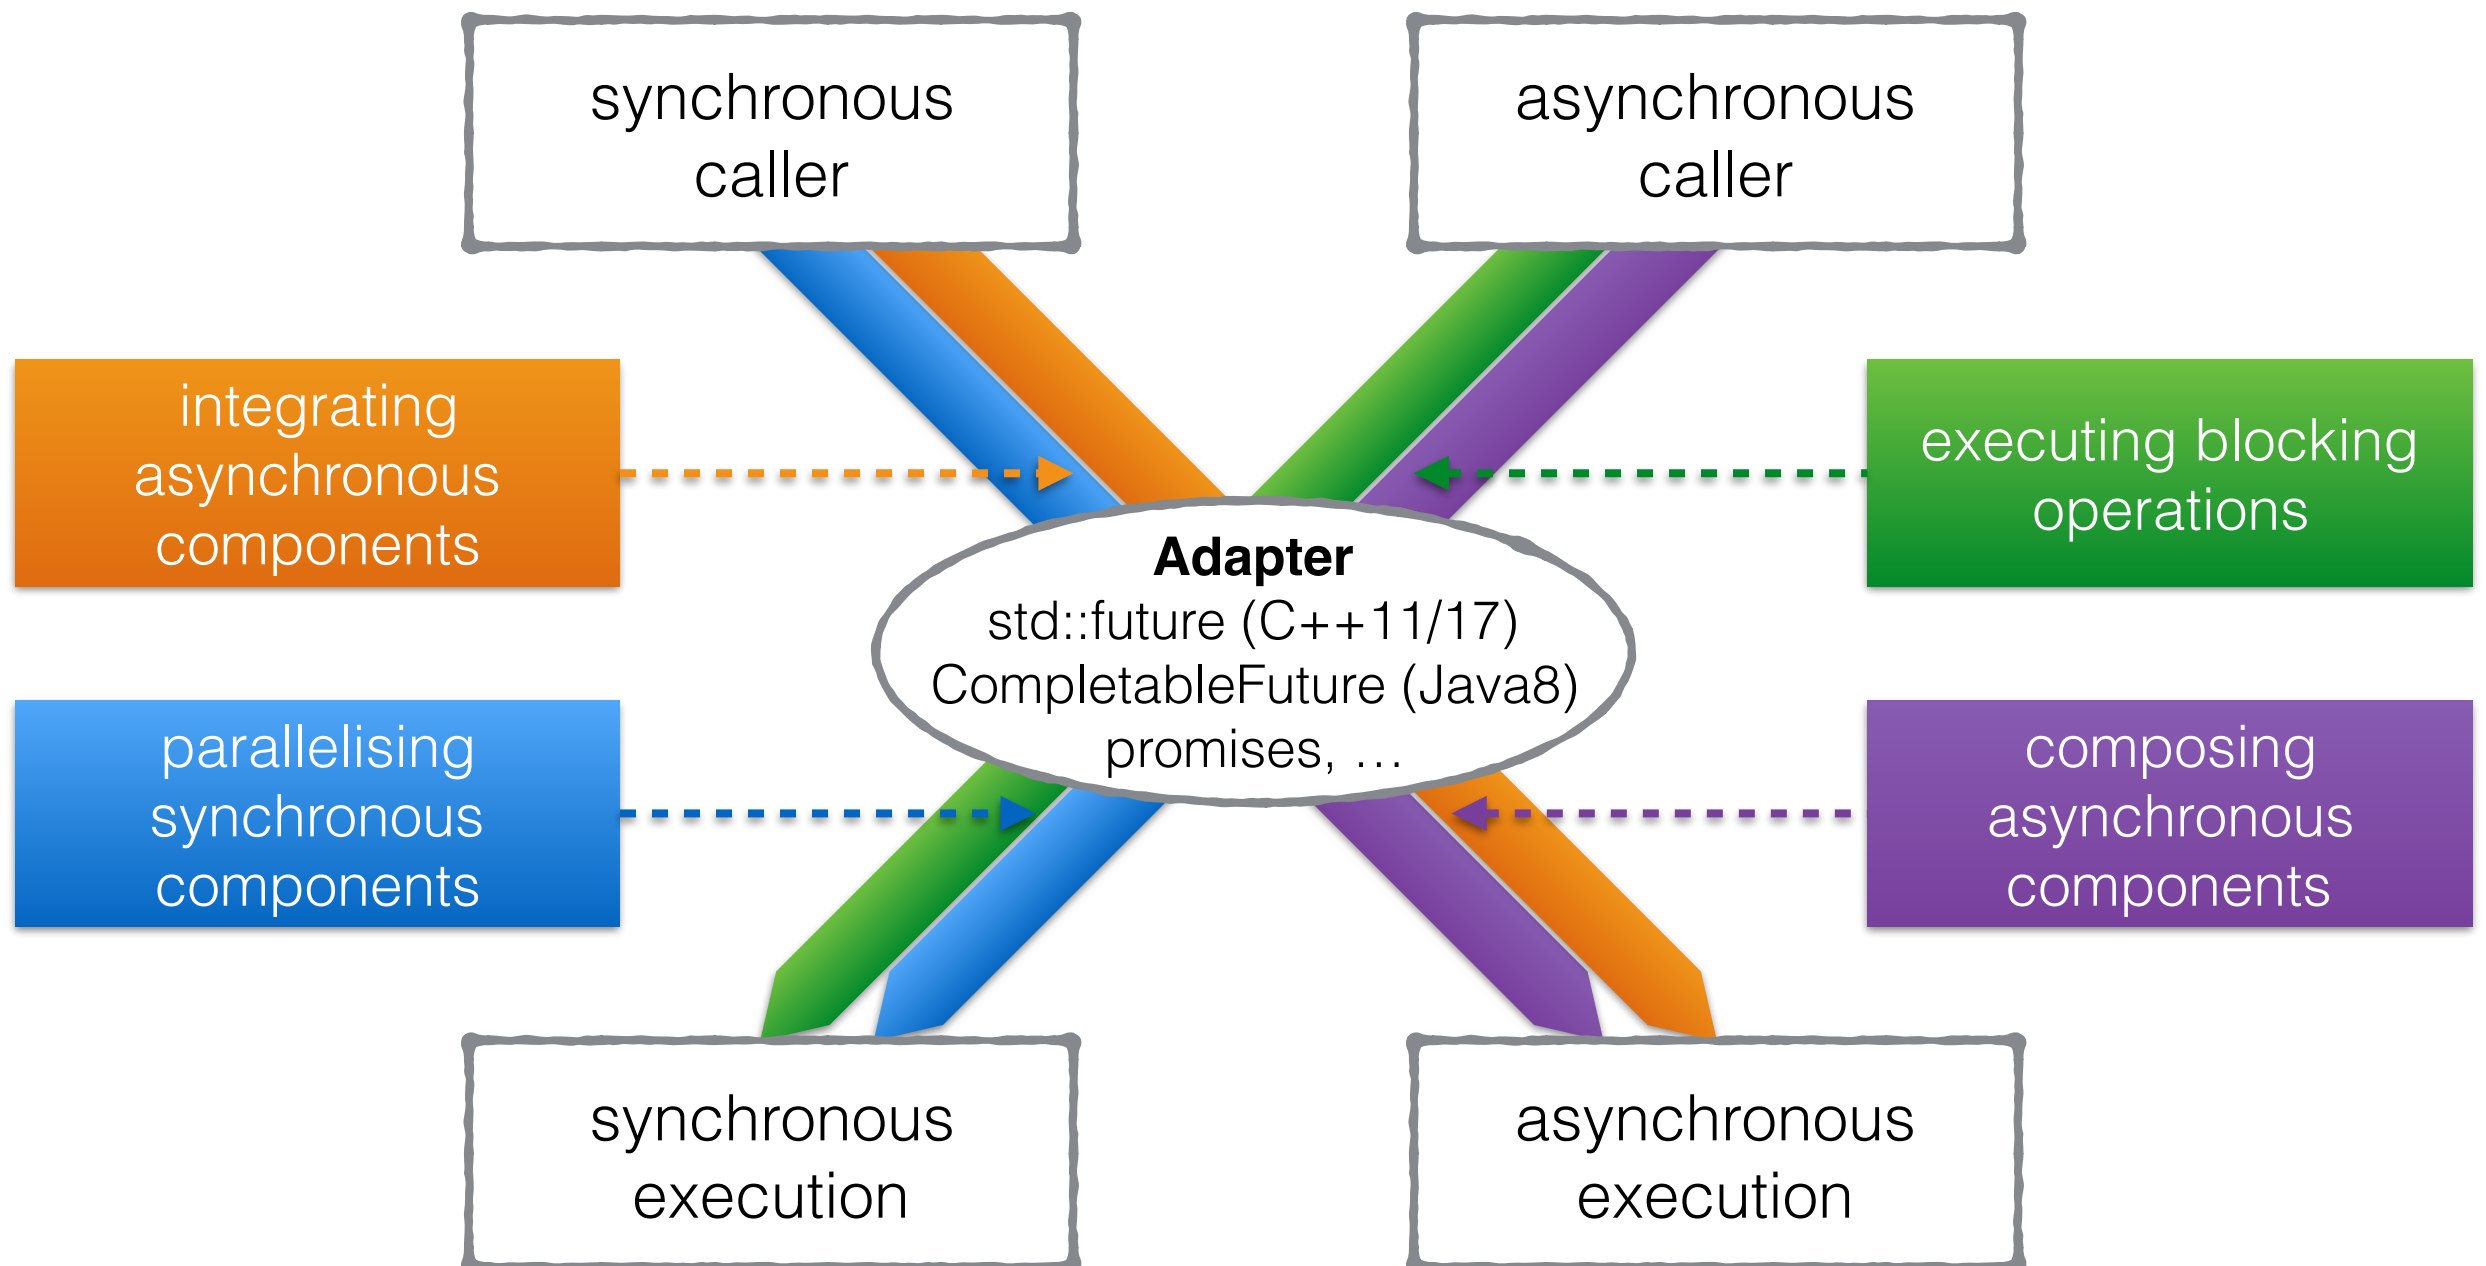
## **third party libraries**

- incomplete, incompatible

# Mitigation



# Interplay



# Summary

- Multi-Core
- RAM and platform peculiarities
- non-sequential flows
- blocking operations
- hybrids





# Sources

- “[The C10K problem](#)” by Dan Kegel
- “[Scaling in the Linux Networking Stack](#)” by Tom Herbert and Willem de Bruijn
- “Comparing the performance of web server architectures.” by Pariag, David, et al.
- “[Hype-Zyklus nach Gartner Inc.](#)” by “ldotter” is licensed under [CC BY-SA 3.0](#)
- “[Arcadia pocket watch 1859](#)” by “Suebun” is licensed under [CC BY-SA 3.0](#)
- “An icon in the Tango! theme style: [Blue person](#). [Green person](#). [Orange person](#). [Purple person](#).” by “Inductiveload” is released into [public domain](#).
- “[Emblem persons blue green](#)” by “Con-struct” is licensed under [CC BY-SA 3.0](#)
- “[View refresh](#)” by “[The people from the Tango! project](#)” is released into [public domain](#).
- “[Korganizer Icon](#)” by “Umut Pulat” is licensed under [GNU Lesser General Public License](#)
- “[Proactor design pattern](#)” by “Christopher M. Kohlhoff” is licensed under [Boost Software License 1.0](#)
- “[AE.svg](#)” by “C.Thure” is released into [public domain](#)
- “[Question mark made of question marks](#)” by “Khaydock” is licensed under [CC BY-SA 3.0](#)

